

HAR_LSTM

July 31, 2019

```
In [1]: # Importing Libraries

In [1]: import pandas as pd
import numpy as np

In [129]: # Activities are the class labels
# It is a 6 class classification
ACTIVITIES = {
    0: 'WALKING',
    1: 'WALKING_UPSTAIRS',
    2: 'WALKING_DOWNSTAIRS',
    3: 'SITTING',
    4: 'STANDING',
    5: 'LAYING',
}

# Utility function to print the confusion matrix
def confusion_matrixn_matrix(Y_true, Y_pred):
    Y_true = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_true, axis=1)])
    Y_pred = pd.Series([ACTIVITIES[y] for y in np.argmax(Y_pred, axis=1)])

    return pd.crosstab(Y_true, Y_pred, rownames=['True'], colnames=['Pred'])
```

0.0.1 Data

```
In [3]: # Data directory
DATADIR = 'UCI_HAR_Dataset'

In [4]: # Raw data signals
# Signals are from Accelerometer and Gyroscope
# The signals are in x,y,z directions
# Sensor signals are filtered to have only body acceleration
# excluding the acceleration due to gravity
# Triaxial acceleration from the accelerometer is total acceleration
SIGNALS = [
    "body_acc_x",
    "body_acc_y",
    "body_acc_z",
```

```

        "body_gyro_x",
        "body_gyro_y",
        "body_gyro_z",
        "total_acc_x",
        "total_acc_y",
        "total_acc_z"
    ]

```

In [5]: *# Utility function to read the data from csv file*

```

def _read_csv(filename):
    return pd.read_csv(filename, delim_whitespace=True, header=None)

# Utility function to load the load
def load_signals(subset):
    signals_data = []

    for signal in SIGNALS:
        filename = f'UCI_HAR_Dataset/{subset}/Inertial Signals/{signal}_{subset}.txt'
        signals_data.append(
            _read_csv(filename).as_matrix()
        )

    # Transpose is used to change the dimensionality of the output,
    # aggregating the signals by combination of sample/timestep.
    # Resultant shape is (7352 train/2947 test samples, 128 timesteps, 9 signals)
    return np.transpose(signals_data, (1, 2, 0))

```

In [6]: `def load_y(subset):`

```

    """
    The objective that we are trying to predict is a integer, from 1 to 6,
    that represents a human activity. We return a binary representation of
    every sample objective as a 6 bits vector using One Hot Encoding
    (https://pandas.pydata.org/pandas-docs/stable/generated/pandas.get\_dummies.html)
    """
    filename = f'UCI_HAR_Dataset/{subset}/y_{subset}.txt'
    y = _read_csv(filename)[0]

    return pd.get_dummies(y).as_matrix()

```

In [7]: `def load_data():`

```

    """
    Obtain the dataset from multiple files.
    Returns: X_train, X_test, y_train, y_test
    """
    X_train, X_test = load_signals('train'), load_signals('test')
    y_train, y_test = load_y('train'), load_y('test')

    return X_train, X_test, y_train, y_test

```

```
In [8]: # Importing tensorflow
        np.random.seed(42)
        import tensorflow as tf
        tf.set_random_seed(42)
```

C:\Users\hp\Anaconda3\lib\site-packages\h5py__init__.py:36: FutureWarning: Conversion of the
from ._conv import register_converters as _register_converters

```
In [9]: # Configuring a session
        session_conf = tf.ConfigProto(
            intra_op_parallelism_threads=1,
            inter_op_parallelism_threads=1
        )
```

```
In [10]: # Import Keras
         from keras import backend as K
         sess = tf.Session(graph=tf.get_default_graph(), config=session_conf)
         K.set_session(sess)
```

Using TensorFlow backend.

```
In [11]: # Importing libraries
         from keras.models import Sequential
         from keras.layers import LSTM
         from keras.layers.core import Dense, Dropout
```

```
In [79]: # Initializing parameters
        """
        epochs = 30
        batch_size = 16
        n_hidden = 64
        drop_out = 0.5

        epochs = 30
        batch_size = 16
        n_hidden = 64
        drop_out = 0.25

        epochs = 30
        batch_size = 32
        n_hidden = 64
        drop_out = 0.25
        """
```

*#this configuration is not improving our accuracy
accuracy is around 90.*

this is also not working well for us; accuracy 91

*#this configuration seems to me as it's doing over
accuracy increase to a certain point and then it*

```

epochs = 30                                # using this configuration we 92.81% almost got 93% accu
batch_size = 32                            # after 27 epochs we are getting higher accuracy 93.72 w
n_hidden = 128                             # almost equal to 94%
drop_out = 0.5

```

```

In [110]: # Utility function to count the number of classes
def _count_classes(y):
    return len(set([tuple(category) for category in y]))

```

```

In [111]: # Loading the train and test data
X_train, X_test, Y_train, Y_test = load_data()

```

```

In [112]: timesteps = len(X_train[0])
input_dim = len(X_train[0][0])
n_classes = _count_classes(Y_train)

```

```

print(timesteps)
print(input_dim)
print(len(X_train))

```

```

128
9
7352

```

- Defining the Architecture of LSTM

```

In [83]: import warnings
warnings.filterwarnings("ignore")
# Initiliazing the sequential model
model = Sequential()
# Configuring the parameters
model.add(LSTM(n_hidden, input_shape=(timesteps, input_dim)))
# Adding a dropout layer
model.add(Dropout(drop_out))
# Adding a dense output layer with sigmoid activation
model.add(Dense(n_classes, activation='sigmoid'))
model.summary()

```

Layer (type)	Output Shape	Param #
lstm_13 (LSTM)	(None, 128)	70656
dropout_13 (Dropout)	(None, 128)	0
dense_13 (Dense)	(None, 6)	774
Total params: 71,430		

Trainable params: 71,430
Non-trainable params: 0

```
-----  
  
In [101]: import warnings  
          warnings.filterwarnings("ignore")  
          # Compiling the model  
          model.compile(loss='categorical_crossentropy',  
                        optimizer='rmsprop',  
                        metrics=['accuracy'])
```

```
In [85]: import warnings  
          warnings.filterwarnings("ignore")  
  
          from datetime import datetime  
          start = datetime.now()  
  
          # Training the model  
          model.fit(X_train,  
                    Y_train,  
                    batch_size=batch_size,  
                    validation_data=(X_test, Y_test),  
                    epochs=epochs)  
  
          print("Time taken : ", datetime.now() - start)
```

Train on 7352 samples, validate on 2947 samples

```
Epoch 1/30  
7352/7352 [=====] - 46s 6ms/step - loss: 1.3462 - acc: 0.4006 - val_loss: 1.3462  
Epoch 2/30  
7352/7352 [=====] - 44s 6ms/step - loss: 1.2504 - acc: 0.4373 - val_loss: 1.2504  
Epoch 3/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.8766 - acc: 0.5884 - val_loss: 0.8766  
Epoch 4/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.7814 - acc: 0.6404 - val_loss: 0.7814  
Epoch 5/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.6919 - acc: 0.6829 - val_loss: 0.6919  
Epoch 6/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.5751 - acc: 0.7703 - val_loss: 0.5751  
Epoch 7/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.4046 - acc: 0.8587 - val_loss: 0.4046  
Epoch 8/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.2885 - acc: 0.8987 - val_loss: 0.2885  
Epoch 9/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.2360 - acc: 0.9146 - val_loss: 0.2360  
Epoch 10/30  
7352/7352 [=====] - 44s 6ms/step - loss: 0.2128 - acc: 0.9233 - val_loss: 0.2128
```

```

Epoch 11/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1994 - acc: 0.9301 - val_loss: 0.1994
Epoch 12/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1666 - acc: 0.9404 - val_loss: 0.1666
Epoch 13/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1761 - acc: 0.9369 - val_loss: 0.1761
Epoch 14/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1522 - acc: 0.9456 - val_loss: 0.1522
Epoch 15/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1502 - acc: 0.9434 - val_loss: 0.1502
Epoch 16/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1492 - acc: 0.9452 - val_loss: 0.1492
Epoch 17/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1485 - acc: 0.9387 - val_loss: 0.1485
Epoch 18/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1575 - acc: 0.9433 - val_loss: 0.1575
Epoch 19/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1465 - acc: 0.9442 - val_loss: 0.1465
Epoch 20/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1386 - acc: 0.9495 - val_loss: 0.1386
Epoch 21/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1373 - acc: 0.9505 - val_loss: 0.1373
Epoch 22/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1227 - acc: 0.9516 - val_loss: 0.1227
Epoch 23/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1495 - acc: 0.9414 - val_loss: 0.1495
Epoch 24/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1212 - acc: 0.9509 - val_loss: 0.1212
Epoch 25/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1266 - acc: 0.9520 - val_loss: 0.1266
Epoch 26/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1345 - acc: 0.9460 - val_loss: 0.1345
Epoch 27/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1306 - acc: 0.9478 - val_loss: 0.1306
Epoch 28/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1619 - acc: 0.9482 - val_loss: 0.1619
Epoch 29/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1267 - acc: 0.9520 - val_loss: 0.1267
Epoch 30/30
7352/7352 [=====] - 44s 6ms/step - loss: 0.1284 - acc: 0.9504 - val_loss: 0.1284
Time taken : 0:22:06.036482

```

```

In [ ]: # Confusion Matrix
        print(confusion_matrix(Y_test, model.predict(X_test)))

In [22]: score = model.evaluate(X_test, Y_test)

2947/2947 [=====] - 1s 461us/step

```

```
In [23]: score
```

```
Out[23]: [0.39372028857254565, 0.9019341703427214]
```

- With a simple 2 layer architecture we got 90.09% accuracy and a loss of 0.30
- We can further improve the performance with Hyperparameter tuning

0.1 LSTM model with 2 layers

```
In [121]: """                                this configuration gives accuracy of 90.84
          epochs_m2 = 30
          batch_size_m2= 32
          n_hidden_layer1 = 128
          n_hidden_layer2 =32
          drop_out_1 = 0.2
          drop_out_2 = 0.5
          """
```

```
epochs_m2 = 30
batch_size_m2= 32
n_hidden_layer1 = 128
n_hidden_layer2 =64
drop_out_1 = 0.2
drop_out_2 = 0.5
```

```
In [122]: # Initiliazing the sequential model
          model2 = Sequential()
          # Configuring the parameters
          model2.add(LSTM(n_hidden_layer1, return_sequences=True, input_shape=(timesteps, input_shape[1])))
          # Adding a dropout layer
          model2.add(Dropout(drop_out_1))

          model2.add(LSTM(n_hidden_layer2))
          # Adding a dropout layer
          model2.add(Dropout(drop_out_2))
          # Adding a dense output layer with sigmoid activation
          model2.add(Dense(n_classes, activation='sigmoid'))
          model2.summary()
```

Layer (type)	Output Shape	Param #
lstm_20 (LSTM)	(None, 128, 128)	70656
dropout_20 (Dropout)	(None, 128, 128)	0
lstm_21 (LSTM)	(None, 64)	49408

```

-----
dropout_21 (Dropout)          (None, 64)          0
-----
dense_18 (Dense)              (None, 6)            390
=====
Total params: 120,454
Trainable params: 120,454
Non-trainable params: 0
-----

```

```

In [123]: # Compiling the model
          model2.compile(loss='categorical_crossentropy',
                        optimizer='rmsprop',
                        metrics=['accuracy'])

```

```

In [124]: import warnings
          warnings.filterwarnings("ignore")

          from datetime import datetime
          start = datetime.now()

          # Training the model
          model2.fit(X_train,
                    Y_train,
                    batch_size=batch_size_m2,
                    validation_data=(X_test, Y_test),
                    epochs=epochs_m2)

          print("Time taken : ", datetime.now() - start)

```

Train on 7352 samples, validate on 2947 samples

```

Epoch 1/30
7352/7352 [=====] - 94s 13ms/step - loss: 1.1846 - acc: 0.4838 - val_
Epoch 2/30
7352/7352 [=====] - 94s 13ms/step - loss: 0.8015 - acc: 0.6364 - val_
Epoch 3/30
7352/7352 [=====] - 92s 12ms/step - loss: 0.6878 - acc: 0.7087 - val_
Epoch 4/30
7352/7352 [=====] - 92s 12ms/step - loss: 0.6659 - acc: 0.7042 - val_
Epoch 5/30
7352/7352 [=====] - 92s 12ms/step - loss: 0.5136 - acc: 0.7734 - val_
Epoch 6/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.3751 - acc: 0.8637 - val_
Epoch 7/30
7352/7352 [=====] - 92s 12ms/step - loss: 0.2722 - acc: 0.9154 - val_
Epoch 8/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.2314 - acc: 0.9242 - val_

```



```

Epoch 9/30
7352/7352 [=====] - 94s 13ms/step - loss: 0.1961 - acc: 0.9348 - val_
Epoch 10/30
7352/7352 [=====] - 105s 14ms/step - loss: 0.1734 - acc: 0.9389 - val_
Epoch 11/30
7352/7352 [=====] - 98s 13ms/step - loss: 0.1789 - acc: 0.9403 - val_
Epoch 12/30
7352/7352 [=====] - 96s 13ms/step - loss: 0.1568 - acc: 0.9416 - val_
Epoch 13/30
7352/7352 [=====] - 99s 14ms/step - loss: 0.1686 - acc: 0.9425 - val_
Epoch 14/30
7352/7352 [=====] - 96s 13ms/step - loss: 0.1430 - acc: 0.9475 - val_
Epoch 15/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1774 - acc: 0.9434 - val_
Epoch 16/30
7352/7352 [=====] - 93s 13ms/step - loss: 0.1591 - acc: 0.9489 - val_
Epoch 17/30
7352/7352 [=====] - 98s 13ms/step - loss: 0.1377 - acc: 0.9494 - val_
Epoch 18/30
7352/7352 [=====] - 99s 13ms/step - loss: 0.1294 - acc: 0.9513 - val_
Epoch 19/30
7352/7352 [=====] - 98s 13ms/step - loss: 0.2095 - acc: 0.9359 - val_
Epoch 20/30
7352/7352 [=====] - 94s 13ms/step - loss: 0.1286 - acc: 0.9479 - val_
Epoch 21/30
7352/7352 [=====] - 93s 13ms/step - loss: 0.1432 - acc: 0.9463 - val_
Epoch 22/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1234 - acc: 0.9531 - val_
Epoch 23/30
7352/7352 [=====] - 94s 13ms/step - loss: 0.1266 - acc: 0.9504 - val_
Epoch 24/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1574 - acc: 0.9478 - val_
Epoch 25/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1226 - acc: 0.9524 - val_
Epoch 26/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1339 - acc: 0.9456 - val_
Epoch 27/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1276 - acc: 0.9532 - val_
Epoch 28/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1432 - acc: 0.9487 - val_
Epoch 29/30
7352/7352 [=====] - 93s 13ms/step - loss: 0.1335 - acc: 0.9489 - val_
Epoch 30/30
7352/7352 [=====] - 92s 13ms/step - loss: 0.1427 - acc: 0.9434 - val_
Time taken : 0:47:09.180834

```

```
In [134]: score = model2.evaluate(X_test, Y_test)
```

```
2947/2947 [=====] - 7s 2ms/step
```

```
In [135]: score
```

```
Out[135]: [0.31484989217882003, 0.9216152019002375]
```

0.2 Conclusion

In classic machine learning models we got accuracy of around 96% and here we also tried to get the accuracy close to 96% by * tuning the number of lstm units * experimenting with drop out value * and by also adding a secnd hidden layer.

For our first model with single hidden layer I tried various configurations :

```
In [142]: from prettytable import PrettyTable
          ##### pretty table for model 1 with 1 lstm layer #####
          number      = [1, 2, 3, 4]
          epochs       = [30, 30, 30, 30]
          batch_size   = [16, 16, 32, 32]
          n_hidden     = [64, 64, 64, 128]
          drop_out     = [0.25, 0.25, 0.25, 0.5]
          accuracy     = [90.6, 91.45, 90.84, 92.81]

          # Initializing prettytable # Adding columns
          ptable = PrettyTable()
          ptable.add_column("Configuration",number)
          ptable.add_column("Epochs", epochs)
          ptable.add_column("Batch Size",batch_size)
          ptable.add_column("Hidden Layer",n_hidden)
          ptable.add_column("Dropout",drop_out)
          ptable.add_column("Accuracy",accuracy)
          #Printing the Table
          print(ptable)
```

Configuration	Epochs	Batch Size	Hidden Layer	Dropout	Accuracy
1	30	16	64	0.25	90.6
2	30	16	64	0.25	91.45
3	30	32	64	0.25	90.84
4	30	32	128	0.5	92.81

```
In [144]: ##### pretty table for model 2 with 2 lstm layesr #####
          number      = [1, 2]
          epochs       = [30, 30]
          batch_size   = [32, 32]
```

```

n_hidden_layer1 = [128, 128]
n_hidden_layer2 = [32, 64]
drop_out_1      = [0.2, 0.5]
drop_out_2      = [0.2, 0.5]
accuracy        = [90.84, 92.16]

# Initializing prettytable # Adding columns
ptable = PrettyTable()
ptable.add_column("Configuration", number)
ptable.add_column("Epochs", epochs)
ptable.add_column("Batch Size", batch_size)
ptable.add_column("Hidden Layer", n_hidden_layer1)
ptable.add_column("Hidden Layer", n_hidden_layer2)
ptable.add_column("Dropout", drop_out_1)
ptable.add_column("Dropout", drop_out_2)
ptable.add_column("Accuracy", accuracy)
#Printing the Table
print(ptable)

```

Configuration	Epochs	Batch Size	Hidden Layer	Hidden Layer	Dropout	Dropout	Accuracy
1	30	32	128	32	0.2	0.2	90.84
2	30	32	128	64	0.5	0.5	92.16

- Note: One observation which I made while training our lstm models is that after training the model more than one time we get differnt values of accuracy i.e for same configuration we might get different accuracy values every time we train it over again.

So the results may improve if we train these models again with same configuration. The final configurations for each model i.e model 1 and model 2 are good and we can get accuracy closer to 96%.